

Final Model Report: Enhancing Chatbot with Hierarchical Indices for RAG and Summarization

1. Analytic Approach

- **Target Definition:**
The primary goal of the project was to create a faithful financial but, by enhancing a baseline chatbot's ability to retrieve and generate relevant, accurate responses. This was achieved by integrating hierarchical indices into the RAG (Retrieval-Augmented Generation) framework and incorporating summarization capabilities.
- **Inputs (Description):**
 - User queries (natural language inputs)
 - Knowledge base documents (structured and unstructured text)
 - Hierarchical indices for improved document retrieval
- **Model Built:**
The solution consists of a fine-tuned LLM (Large Language Model) integrated with a hierarchical indexing mechanism to optimize retrieval processes, coupled with summarization modules to distill information effectively.

2. Solution Description

- **Simple Solution Architecture:**
Description and Implementation:
The solution consists of changes and improvements to baseline functionality inside the streaming and inference pipeline:
 - **Streaming Pipeline:**
After extracting and preprocessing the documents as in the baseline, and before inserting them into Qdrant, we added the following:
 - For each document, using the ChatGPT API, we classify it into:
 - **Sector:** The broad industry/field to which the document belongs (e.g., Healthcare, Technology).
 - **Subject:** The specific company/person/other subject mentioned in the document (e.g., Microsoft under the Technology sector).
 - **Event Type:** The type of event/activity described (e.g., for Microsoft under Technology, an event type might be a product launch).

- The classification is performed in separate steps:
 - Query to identify the **Sector**.
 - Query to identify the **Subject** within the sector.
 - Query to identify the **Event Type** within the subject and sector.
 - The resulting triplet (Sector-Subject-Event Type) is saved into the document's payload as the collection name.
 - The document is then inserted into Qdrant with its new collection name attribute.
 - While querying for documents, we maintain a tree-like structure inside a JSON file to store the hierarchy. For each document, we provide existing options for each level, asking GPT to choose the correct option if it exists or invent one otherwise.
 - Note: All documents are still saved inside the same Qdrant collection (like in the baseline). The collection name is only an attribute inside each document's payload.
- **Outcomes of Changes to the Streaming Pipeline:**
 - All documents are stored in Qdrant with a new field attached to their payload as the collection name.
 - The fully-created hierarchy is stored inside our JSON file.
- **Inference Pipeline:**
 Our changes focus on the context-retrieval part:
 - Similarly to the streaming pipeline, we classify the user query into Sector, Subject, and Event Type using GPT — this time **ONLY** allowing options that exist inside the JSON for each level. However, we allow several subjects and event types to capture broader context at most 3 subjects and at most 5 event types under each subject.
 - For all resulting hierarchy triplets, we fetch documents from Qdrant and select the top 10 most similar results to the query.
 - If a chosen document lacks the generated summary, we query GPT to create a concise summary.
 - Finally, we concatenate all summaries and send them to the Falcon 7B model as context for the user's query.

- **Outcomes of Changes to the Inference Pipeline:**
 - The chatbot retrieve and generate relevant, accurate responses
 - Less hallucinations and answers are more grounded to truth

3. Data

- **Source:**
Alpaca News API for batch and online learning datasets for fine-tuning, user query logs for evaluation.
- **Data Schema:**
- **In the streaming pipeline:**
 - **Alpaca news Document:**

```
{
  "type": "object",
  "properties": {
    "T": { "type": "string" },
    "id": { "type": "integer" },
    "headline": { "type": "string" },
    "summary": { "type": "string" },
    "author": { "type": "string" },
    "created_at": { "type": "string", "format": "date-time" },
    "updated_at": { "type": "string", "format": "date-time" },
    "url": { "type": "string", "format": "uri" },
    "content": { "type": "string" },
    "symbols": {
      "type": "array",
      "items": { "type": "string" }
    }
  },
}
```

```

        "source": { "type": "string" }
    },
    "required": ["T", "id", "headline", "created_at", "url", "content",
                 "symbols", "source"]
}

```

- **Data in the training pipeline:**

- Q&A object:

```

{
    "type": "object",
    "properties": {
        "about_me": { "type": "string" },
        "context": { "type": "string" },
        "response": { "type": "string" }
    },
    "required": ["about_me", "context", "response"]
}

```

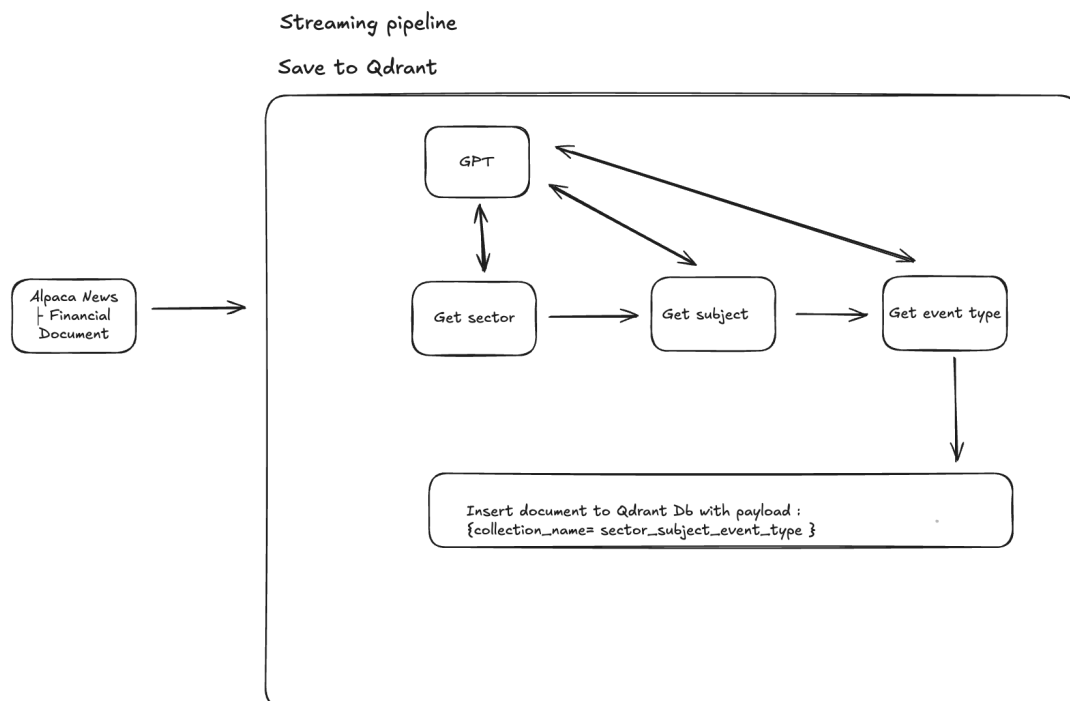
Data in the inference pipeline:

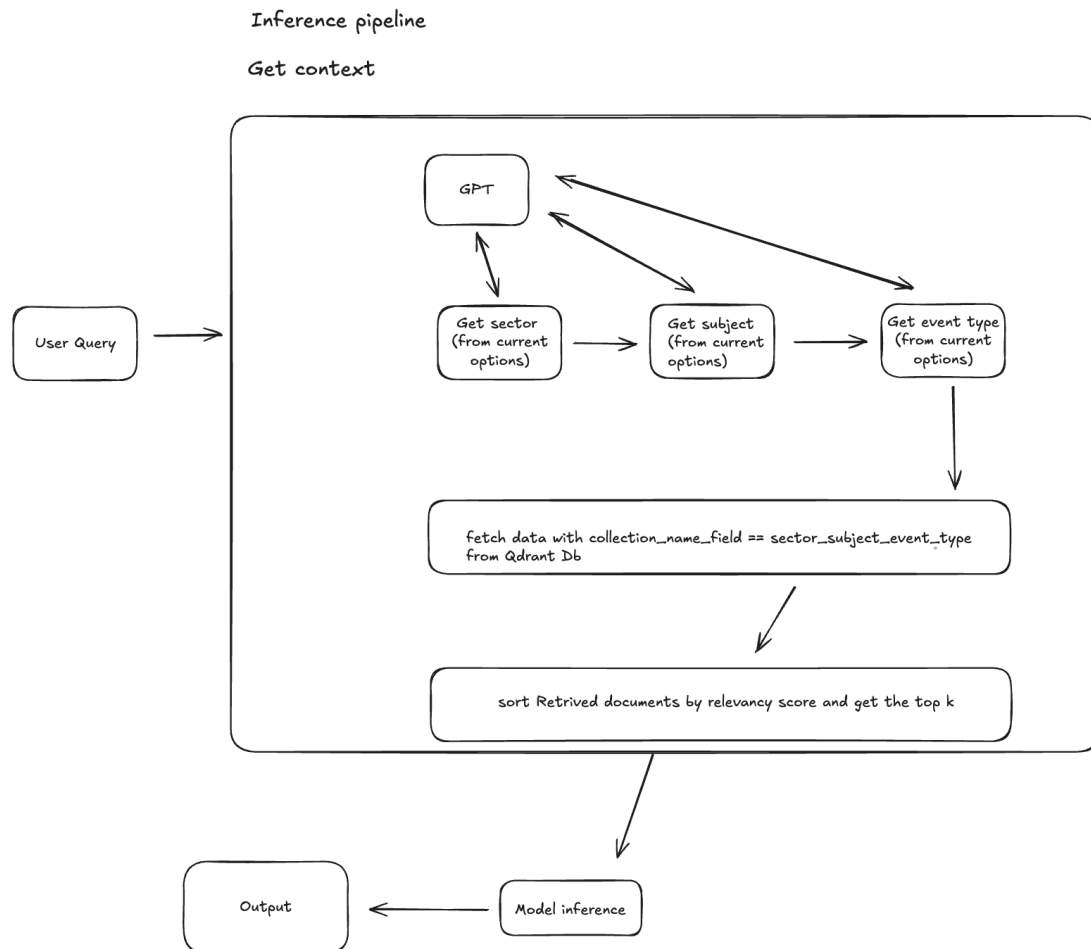
- User Query:
 - “ Query Text”:{ "type": "string" }
- **Selection (Dates, Segments):**
- **Streaming pipeline:**
when running the streaming pipeline using command :”run batch” it takes data from alpaca from the past 8 days, and when working with “run batch dev” it takes from the last 2 days
- **Training pipeline:**
- The model was trained on data from 01.01.2023-05.01.2023 .

4. Features

- **List of Raw and Derived Features:**
 - Raw: Document text, document summary, query text, metadata
 - Derived: TF-IDF vectors, embeddings from pre-trained models, hierarchical index scores, sector-subject-event type classification
- **Importance Ranking:**
 - Hierarchical index scores
 - Embedding similarity
 - Query-document relevance scores
 - Sector-Subject-Event Type classification

5. Algorithm





- **Learners Used:**

- Falcon 7B - Pre-trained LLM from hugging face, after finetuning on financial data questions.

6. Results

- **Performance Metrics:**

- **Answer similarity:** 0.52 -> 0.64 (+0.12)
- **Faithfulness :** 0.2932 -> 0.5392 (+0.2459)